

Sarah Connor ■

AI Security Engineer · Adversarial ML & Red Teaming

✉ sarah@sconnor.dev · 📍 Austin, TX (US)

🐙 github.com/sconnor · 🔗 linkedin.com/in/sarahconnor

> 1 September 2026

Dear Hiring Manager,

I am writing to apply for the Staff AI Security Engineer role on your detection team. Over the past six years I have moved from applied machine learning into security engineering and, most recently, into building the adversarial and red-team defenses that keep production LLM systems safe to ship.

At Cyberdyne Systems I built an automated prompt-injection fuzzing harness that surfaced more than 40 high-severity jailbreaks across six production LLM features before they ever reached customers, and led adversarial robustness reviews that cut a fraud-detection model's evasion rate from 18% to under 2%. I also designed the model-extraction detection pipeline now running across every customer-facing inference endpoint we operate.

Earlier, at Skynet Analytics, I authored the organization's first threat model for LLM-integrated services — since adopted company-wide — and embedded security gating directly into the MLOps CI/CD pipeline, cutting the time to patch critical model-serving CVEs from three weeks to four days.

Your team's work on production LLM safety is exactly the kind of problem I want to keep solving, and I would welcome the chance to bring that experience to your organization.

Thank you for your consideration.

Best regards,
Sarah Connor